AIMS | African Institute for Mathematical Sciences
IN AFRICA'S YOUTH, THE FUTURE OF SCIENCE

AIMS KTT FELLOWSHIP

Hackathon 2026

Candidate Briefing

Everything you need to know about the 3-day evaluation: what to submit, how you'll be scored, and how to prepare for the final pitch.

22 – 24 April 2026 · Virtual · 10:00–19:00 CAT

What this hackathon is — and what it is not

The KTT Fellowship Hackathon selects 9 fellows from a pool of 52 applicants.

Over three days you will be evaluated on a single challenge you receive. The challenge asks you to build something useful for African realities — low bandwidth, multilingual users, constrained devices — and to defend your work video and/or live.

It is not a coding marathon. It is a four-hour focused build followed by structured review. We value clear thinking, honest trade-offs, and real user empathy over flashy demos.

52 → 9

A competitive funnel across three days of independent review.

4-hour slot

You will be given one challenge. You have four hours to build, then submit your repo, model, data, and a 4-min video.

Live defense

On Days 3, you may be asked to open your repo on a screen-share and answer questions. This is how we confirm the work is yours.

How 52 become 9

52 candidates submit

All submissions reviewed by 5 × 2 evaluators

27 advance to Day 2

Collaborative review by 3 × 3 evaluators

15 advance to Day 3

Senior panel of 4 judges, full consensus

9 Fellows selected

Announced at the close of Day 3

Groups grow in size and seniority as the pool shrinks — the closer you get to the top, the more experienced the panel.

Three days, one timezone (CAT)

Day 1
Tue 22 Apr 2026

Initial Screening

16:00 – 18:00 CAT

Your submissions are reviewed. No live session required from you — be available for clarification.

Day 2
Wed 23 Apr 2026

Intermediate Selection

16:00 – 18:00 CAT

If you advanced: a 3-evaluator panel.

Day 3
Thu 24 Apr 2026

Final Selection & Pitch

15:00 – 18:00 CAT

Finalists pitch to the senior panel. Fellows announced at the close of day.

What goes in your submission package

1 4-Minute Video

Required

Walkthrough of your solution with live code + live demo. Structure is prescribed in your challenge brief.

2 Code (GitHub or GitLab)

Required

Public repo. README that runs in ≤ 2 commands on free Colab CPU. LICENSE file.

3 Model (Hugging Face / Ollama)

Required

Hosted model with a model card. For unusual formats, Google Drive or OneDrive is acceptable with a direct download link.

4 Dataset or generator

Required

Training/test data or a reproducible generator script. For > 100 MB, host on HF Datasets, Kaggle, or Drive and link from README.

+ Live demo

Optional

Interactive demo of your solution — deployed endpoint, Gradio, or similar.

5 process_log.md + SIGNED.md

Required

Your declared LLM / tool use and the honor code signed at the root of your repo.

Six criteria, a 1–5 scale, transparent weights

Technical Quality & Code — **20%**

Scored on: GitHub repo

Model / Algorithm Performance — **20%**

Scored on: Reported metrics

Data Handling & Methodology — **15%**

Scored on: Notebook + dataset

Product & Business Adaptation — **20%**

Scored on: Required artifact
KTT signature criterion

Communication & Documentation — **15%**

Scored on: README + 4-min video

Innovation & Problem-Solving — **10%**

Scored on: Overall submission

Each criterion is scored 1–5. Missing deliverables score 1. Evaluators use the full scale — don't assume safety at 3.

5	4	3	2	1
Excellent	Very Good	Good	Fair	Poor

The asynchronous artifact that matters most

Structure your video exactly as below. Your challenge brief lists the exact function to walk through, the demo input to run live, and three spoken questions you must answer.

0:00 – 0:30	0:30 – 1:30	1:30 – 2:30	2:30 – 3:30	3:30 – 4:00
On-camera intro	**Live code walk**	**Live demo**	**Product & Business**	**Three answers**
Name, challenge ID, the hardest trade-off you made.	Screen-share. Open the named function in your own repo.	Run the specified input. Terminal output must be visible.	Walk your adaptation artifact on screen.	Technical, product, local-context. Spoken, unscripted.

✓ DO

Keep the camera on for intro + outro
Show real terminal output, not screenshots
Speak unscripted when answering the 3 Qs

✗ AVOID

Narration over static slides with no code
Pre-rendered outputs with no live execution
Reading a polished LLM-written script verbatim

Free, public, reproducible — pick the right platform per artefact

ARTEFACT	WHERE AND HOW
Code	GitHub or GitLab (public, with README + LICENSE)
Model / checkpoint	Hugging Face Hub (preferred) · Ollama · Google Drive · OneDrive
Dataset / generator	Generator in-repo · Hugging Face Datasets · Kaggle · Drive
4-minute video	YouTube (unlisted) · Vimeo · Google Drive — URL in README
process_log.md	Root of your repo
SIGNED.md (honor code)	Root of your repo — name + date + verbatim text

Rule of thumb: your evaluator should reproduce your work from the README in under two commands on a free Colab CPU.

Wednesday 22 April · 52 → 27

WHAT HAPPENS

Your submission has already been delivered.

Evaluators are organised into 5 groups of 2. Each group processes about 10–11 submissions in a 2-hour window (16:00 – 18:00 CAT).

The approach here is breadth over depth. Your reviewers skim all deliverables, apply the rubric, and advance the top 27.

Decisions are communicated by the end of Day 1.

WHAT WE LOOK FOR AT DAY 1

- All 4 required deliverables present — missing items score 1.
- README that actually runs (not a wishlist).
- The 4-minute video follows the prescribed structure.
- Your Product & Business adaptation has real numbers, users, and trade-offs — not vague paragraphs.
- Your code is honest about what works and what is a known limitation.

Pro tip: if you advance to Day 2, you will be notified the same evening.

Thursday 23 April · 27 → 15

WHAT HAPPENS

3 groups of 3 evaluators review 9 candidates each.

You will share your screen, open your repo, and answer a small number of questions on your own code on video. Expect one ad-hoc modification request (e.g. 'add a unit test for an empty input').

Advancement to Day 3 is announced at the close of Day 2.

HOW TO PREPARE

- Be ready to navigate to a named function without grep.
- Have your process_log.md open — you may be asked about one of your declared prompts.
- Join on a stable connection;
- If you don't know an answer, say so. Honesty is scored positively.

Pro tip: the live defense is where LLM-generated solutions fail most visibly. If you understand your own work, you will do well here.

Friday 24 April · 15 → 9 Fellows

THE PITCH — YOUR 10-MINUTE SLOT WITH THE SENIOR PANEL

0:00 – 4:00

Pitch

You speak — problem, solution, impact, one crisp demo.

4:00 – 7:00

Panel Q&A

3-4 senior judges ask targeted questions.

7:00 – 9:30

Live defense

One modify-on-the-fly in your own repo.

9:30 – 10:00

Close

Your 30-second 'if you fund me, I will ship X by month 3'.

WHAT TO BRING TO YOUR PITCH

- A working demo you can run live in under 60 seconds.
- One memorable number — cost, latency, impact, reach.
- A clear 'next 90 days' plan if selected.

WHAT THE PANEL WILL PROBE

- Judgement under constraint — 'why this trade-off?'
- Empathy for low-tech users — 'show me the grandmother flow'.
- Honesty about limits — 'what breaks your system first?'

The 9 Fellows are announced at the close of Day 3, after the senior panel reaches full consensus.

The single strongest signal — on screen-share, in your voice

Live Defense is how we confirm the work is yours. No LLM, no teammate, no pre-recorded demo. You open your repo, we ask things

Day 1	—	No live session needed. Your submission is reviewed asynchronously.
Day 2	**10–15 min**	One file walk-through, one ad-hoc question, one small modification.
Day 3	**~10 min**	Part of your pitch slot. Code archaeology + modify-on-the-fly + one ethics probe.

A failed Live Defense — cannot explain your own code, cannot modify it — caps your total score. Everything else in the submission becomes secondary.

Before · During · After

BEFORE

From now to the start of your 4-hour slot

- Read your challenge brief end-to-end (including the Video Requirements).
- Prepare your repo: empty GitHub/GitLab, README skeleton, LICENSE.
- Prepare a clean Colab notebook.
- Sign SIGNED.md with the honor code.

DURING

Your 4-hour build slot

- Ship a working baseline within the first 90 minutes, even if simple.
- Log your LLM / tool use in process_log.md as you go (not at the end).
- Leave at least 45 minutes for the video + README polish.
- Stop coding 15 minutes before the end to review.

AFTER

Submission → Day 1 → Day 2 → Day 3

- Submit all links by the deadline. Verify they open in an incognito browser.
- Be available on Day 2 (16h–18h CAT) for judges can request clarification.
- If you make it to Day 3, prepare a 60-second live demo + a 30-second 'next 90 days' pitch.

You may use any tool — own the result

DO — AND DECLARE IT

- Use Claude, ChatGPT, Copilot, Cursor, any other tool. We assume you will.
- Log each tool in process_log.md: name, version, purpose, 3 sample prompts, and 1 prompt you discarded.
- Write a one-paragraph reflection on the single hardest decision you made.
- Sign SIGNED.md at the root of your repo — name, date, honor code verbatim.

DON'T

- Don't have another human do your work.
- Don't hide your tool use — undeclared LLM use is disqualifying.
- Don't paste LLM output you cannot explain. You will be asked to walk specific functions live.
- Don't over-scope. A clean, simple, correct baseline always beats a half-working 'production' solution.

Honor code (to sign in SIGNED.md): "I will declare every LLM or assistant tool I use in my process_log.md. I will not have another human do my work. I will defend my own code in Live Defense."

Things that quietly kill submissions

✗ Over-scoping

Half-shipped features rank below a complete minimal baseline. Choose the smallest version that still solves the user's problem.

✗ README debt

A README that says 'run main.py' but main.py crashes on import will score 1 on Technical Quality. Test your README on a fresh Colab before submitting.

✗ No local context

A Product & Business artefact that says 'deliver via SMS' without a message template, a cost, or a user role loses the KTT signature criterion.

✗ Pretty slides, no code

A polished 4-minute narration with no live terminal and no live code is scored as a structural failure, regardless of visual quality.

✗ Dodged questions

Pretending to know when you don't is heavily penalised at Live Defense. Say 'I didn't get to that' and describe what you would do.

✗ Forgetting the video timing

Under 3 minutes → content missing. Over 4:30 → edit it down. Practice once with a timer.

AIMS | African Institute for Mathematical Sciences

IN AFRICA'S YOUTH, THE FUTURE OF SCIENCE

We built this hackathon to find the engineers who solve the right problem the right way for the right user.

You already submitted. Now show us — through your repo, your video, and your live defense — that the work is yours and that it will travel to the people who need it.

Good luck. We'll see you on Day 3.

Questions? Reach out to the KTT programme team before your 4-hour slot begins.